


Tutorial for A Hitchhiker Guide to Empirical Macro Models¹

Filippo Ferroni (Chicago Fed)

Fabio Canova (Norwegian Business School, CAMP and CEPR)

March 24, 2022

¹The views in this paper are solely the responsibility of the authors and should not be interpreted as reflecting the views of the Federal Reserve Bank of Chicago or any other person associated with the Federal Reserve System. 

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Introduction

- The Empirical Macro Toolbox (EMT) is a package which uses MATLAB (Octave) functions and routines to estimate a number of reduced form and structural macro models with either Classical or Bayesian methods and to forecast.

Introduction

- The Empirical Macro Toolbox (EMT) is a package which uses MATLAB (Octave) functions and routines to estimate a number of reduced form and structural macro models with either Classical or Bayesian methods and to forecast.
- It allows the computation of trend and cycle decompositions, to date turning points and compute statistics of cyclical phases.

Introduction

- The Empirical Macro Toolbox (EMT) is a package which uses MATLAB (Octave) functions and routines to estimate a number of reduced form and structural macro models with either Classical or Bayesian methods and to forecast.
- It allows the computation of trend and cycle decompositions, to date turning points and compute statistics of cyclical phases.
- It allows for inference on parameters; point/density forecasts; measurement of dynamic propagation of *identified* shock; nowcasts and mixed-frequency estimation, system reduction, and regularization in estimation of large datasets.

Introduction

- The Empirical Macro Toolbox (EMT) is a package which uses MATLAB (Octave) functions and routines to estimate a number of reduced form and structural macro models with either Classical or Bayesian methods and to forecast.
- It allows the computation of trend and cycle decompositions, to date turning points and compute statistics of cyclical phases.
- It allows for inference on parameters; point/density forecasts; measurement of dynamic propagation of *identified* shock; nowcasts and mixed-frequency estimation, system reduction, and regularization in estimation of large datasets.
- Source codes with examples can be forked/downloaded at https://github.com/naffe15/BVAR_ or <https://sites.google.com/view/fabio-canova-homepage/home/empirical-macro-toolbox>

Introduction

Network analysis
included! (not optimized
but easy to use)

Regime switching and
higher moments not (yet)
included.

- The Empirical Macro Toolbox (EMT) is a package which uses MATLAB (Octave) functions and routines to estimate a number of reduced form and structural macro models with either Classical or Bayesian methods and to forecast.
- It allows the computation of trend and cycle decompositions, to date turning points and compute statistics of cyclical phases.
- It allows for inference on parameters; point/density forecasts; measurement of dynamic propagation of *identified* shock; nowcasts and mixed-frequency estimation, system reduction, and regularization in estimation of large datasets.
- Source codes with examples can be forked/downloaded at https://github.com/nafe15/BVAR_ or <https://sites.google.com/view/fabio-canova-homepage/home/empirical-macro-toolbox>
- The hitchhiker guide can be downloaded at <https://www.filippoferroni.com/empiricalmacrotoolbox> or <https://sites.google.com/view/fabio-canova-homepage/home/empirical-macro-toolbox>

Bird's eye view of the content

- VAR methods:

Classic and Bayesian estimation, Point/density forecasts, many identification schemes for SVARs, computation of IRF, historical and variance decompositions, missing values/mixed frequency estimation, VARXs, panels of VARs, time varying coefficient and heteroschedastic models.

Bird's eye view of the content

- VAR methods:
Classic and Bayesian estimation, Point/density forecasts, many identification schemes for SVARs, computation of IRF, historical and variance decompositions, missing values/mixed frequency estimation, VARXs, panels of VARs, time varying coefficient and heteroschedastic models.
- Regularization Methods. Network analysis
Lasso, Ridge, Elastic-Net, Bayesian shrinkage estimation. Connectedness and spillovers measures.

Bird's eye view of the content

- VAR methods:
Classic and Bayesian estimation, Point/density forecasts, many identification schemes for SVARs, computation of IRF, historical and variance decompositions, missing values/mixed frequency estimation, VARXs, panels of VARs, time varying coefficient and heteroschedastic models.
- Regularization Methods. Network analysis
Lasso, Ridge, Elastic-Net, Bayesian shrinkage estimation. Connectedness and spillovers measures.
- Compression methods:
Static (PC) and Dynamic Factor models; FAVARs.

Bird's eye view of the content

- VAR methods:
Classic and Bayesian estimation, Point/density forecasts, many identification schemes for SVARs, computation of IRF, historical and variance decompositions, missing values/mixed frequency estimation, VARXs, panels of VARs, time varying coefficient and heteroschedastic models.
- Regularization Methods. Network analysis
Lasso, Ridge, Elastic-Net, Bayesian shrinkage estimation. Connectedness and spillovers measures.
- Compression methods:
Static (PC) and Dynamic Factor models; FAVARs.
- Local projection methods:
Classic and Bayesian estimation, point and density forecasts, structural IRFs.

Bird's eye view of the content

- VAR methods:
Classic and Bayesian estimation, Point/density forecasts, many identification schemes for SVARs, computation of IRF, historical and variance decompositions, missing values/mixed frequency estimation, VARXs, panels of VARs, time varying coefficient and heteroschedastic models.
- Regularization Methods. **Network analysis**
Lasso, Ridge, Elastic-Net, Bayesian shrinkage estimation. Connectedness and spillovers measures.
- Compression methods:
Static (PC) and Dynamic Factor models; FAVARs.
- **Local projection methods:**
Classic and Bayesian estimation, point and density forecasts, structural IRFs.
- Filtering methods:
Extracting trend/cycles and gaps, Dating BC, computing turning points.

Road map I

Topics (with examples in the tutorials as blueprints):

- I Dating turning points. Computation of cyclical statistics.
[USERLOCATION/BVAR/examples/Trend-Cycle-Dating tutorial/example_2_dating.m](#)
- II Trend and cycle decompositions. Comparison of the financial and the real cycles.
[USERLOCATION/BVAR/examples/Trend-Cycle-Dating tutorial/example_1_cycles.m](#)
- III Classical VAR inference (Bayesian estimation with flat prior).
[USERLOCATION/BVAR/examples/BVAR tutorial/example_1_classical.m](#)
- IV VARX: estimation and inference.
[USERLOCATION/BVAR/examples/BVAR tutorial/example_6_VARX.m](#)
- V Bayesian VAR inference: Minnesota and conjugate priors. Chosing Minnesota hyper-parameters.
[USERLOCATION/BVAR/examples/BVAR tutorial/example_2_minn.m](#)
- VI Structural IRFs. Historical and Variance Decompositions.
[USERLOCATION/BVAR/examples/BVAR tutorial/example_3_irf.m](#)
[USERLOCATION/BVAR/examples/BVAR tutorial/example_10_VAR_heterosk.m](#)

Road map II

Topics (examples in the tutorial as blueprints):

VII Regularization methods and network analysis.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_11_connectedness.m`

VIII PCVAR (FAVAR) estimation and inference.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_5_favar.m`

IX DFM estimation and inference.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_12_bdfm.m`

X Panels of VARs.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_8_panels.m`

XI Nowcasts and Mixed-Frequency VAR estimation.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_4_mfvar.m`

XII Point and Density Forecasts. Conditional Forecasts in VARs.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_9_prediction.m`

XIII Local projections and direct forecasts.

`USERLOCATION/BVAR/examples/BVAR tutorial/example_7_LP.m`

Hitting the road

- I Open MATLAB, and add the toolbox to the MATLAB path.

If EMT is stored in C:\USERLOCATION, type in the command window (in your script):

```
addpath C:/USERLOCATION/BVAR/bvartools
```

You need to this every time you start MATLAB. Alternatively, go to the set-path icon in the MATLAB HOME window, click on it, add with BVAR with subfolders, and save.

- II Use the same procedure to install Chris Sims' optimization routines (used to optimally select the Minnesota prior hyper-parameters). In this case, type in the command window

```
addpath C:/USERLOCATION/BVAR/cmintools
```

- III Load the data

- If data is matlab format, type in the command window:

```
load data
```

and check the content in the workspace window.

- If the data comes in a spreadsheet use (standard MATLAB) command:

```
y=xlsread('filename','sheetname')
```

where `filename` is the name of the file, and `sheetname` is the name of the sheet where you have your data (in case you have many sheets).

- For more information about loading data in other formats, type in the command window:

```
help load
```

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Turning point computations I

- Follows Bry and Boschen algorithm as implemented by Harding and Pagan (2006).
- Activated using the command:

```
[dt_] = date_(y, time, freq, tstart, tend, options);
```

where `y` is the data, `time` is the time span (assumed of same length as `y`), `freq` is a string with the frequency of the data; that is, `'m'` (`'q'`) for monthly (quarterly) data; `tstart` (`tend`) is a scalar with the year and the month/quarter with the start (end) of the exercise, e.g. `tstart = [1970 3]` and `tend = [2020 1]`. The final input, `options`, allows customization:

- `options.phase` sets the minimum duration of the estimated phase; default for quarterly (monthly) data is 2 (6).
- `options.cycle` sets the minimum length of the estimated cycle; default for quarterly (monthly) data is 5 (15).
- `options.thresh` bypasses phase and cycle restrictions if peak to trough is larger than a threshold; default is 10.4.
- `options.complete`. When equals one, it uses only complete cycles; else incomplete cycles are used (excess still computed on complete cycles). Default value equals one.
- `options.nrep` is used when multiple datasets (e.g. simulated from a model) are considered. When `options.nrep` is employed, the dating exercise is computed using `options.nrep` \times `T` observations. Default is one.

Turning point computations II

- The output [dt_] is a field structure containing a number of objects:
 - dt_.st is a $T \times 1$ a binary array containing the phases, i.e. expansion (1) and contraction (0).
 - dt_.trinary is a $T \times 1$ trinary array which contains the cycle states, i.e. peak (1), trough (-1) or none of the two (0).
 - dt_.dura is a 2×1 array with the average duration of contractions and of expansions, respectively.
 - dt_.ampl is a 2×1 array with the average amplitude of contractions and of expansions, respectively.
 - dt_.cumm is a 2×1 array with the average cumulative change in contractions and expansions, respectively.
 - dt_.excc is a 2×1 array reporting the percent of excess movements of the triangle area for contraction and expansions, respectively.
 - dt_.durcv is a 2×1 array the coefficient of variation (the inverse of the standard deviation) of the duration of contractions and expansions.
 - dt_.amplcv is a 2×1 array the coefficient of variation of the amplitude of contractions and expansions.
 - dt_.exccv is a 2×1 array the coefficient of variation of the excess movements in contractions and expansions.
- The concordance of turning points (or of cyclical phases) is not automatically computed. See manual (or tutorial examples) on how to do it.
- If y includes more than one series, one needs to loop the date_ function appropriately.

Outline

The trend does not necessarily equal to optimal (state of the economy), which cannot be measured from the data.

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions**
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Trend and Cycle decompositions I

Many decompositions are allowed. Most are univariate; some are bivariate (multivariate).

- **Polynomial trends.** The decomposition is activated using:

```
[dc,dt] = polydet(y,ord,gg, timeplot, varnames);
```

where y is the data, ord is the order of the polynomial (the upper limit is 4) and gg an indicator; if 1 the data, the trend, and the cycle are plotted; $timeplot$ is the time axis; $varnames$ the titles of the plots. The outputs are the estimated cycle dc and the estimated trend dt . If y is a $N \times 1$ vector, you will have N plots on the screen.

- **Polynomial breaking trends.** The decomposition is activated using:

```
[dc,dt] = polydet_break(y,ord,tt,gg, timeplot, varnames);
```

where y is the data, ord is the order of the polynomial (upper limit 4), tt is the break date and gg an indicator; if 1 the data, the trend, and the cycle are plotted. $timeplot$ and $varnames$ indicate the time axis and the titles of the plots. The outputs are the estimated cycle dc and the estimated trend dt . The break date must be selected in advance.

- **Differencing.** No option is set, but can be easily implemented with MATLAB commands (see manual or tutorial example).

Trend and Cycle decompositions II

- **Hamilton (local projection) approach.** The decomposition is activated using:

```
[dc,dt] = hamfilter(y,h,lags,ff,gg, timeplot, varnames);
```


where y is the data, h is the horizon of the projection (upper limit $h = T - \text{lags} - 1$); lags is the number of lags in the projection equation; ff an indicator; if it is equal to 1 a constant is included in the projection; gg an indicator; if 1 the data, the trend, and the cycle are plotted; timeplot and varnames indicate the time axis and the titles of the plots. The outputs are the estimated cycle dc and the estimated trend dt . y could be a $T \times n$ matrix, where n the number of series. Conditioning variables other than the lags of y , are not currently allowed.

- **Unobservable component (UC) 1:** decomposition with random walk trend and AR(2) cycle. Activated by:

```
[dt,dc] = uc1_(y,opts);
```

where y is the data and opts is a field structure that allows customization:

- `opts.figg`: when equals 1 it plots the series, the trend, and the cycle.
- `opts.nsim`: sets the number of MCMC simulations (default 5000).
- `opts.burnin`: sets the number of burn-in draws (default, 30000).
- `opts.rhoyes` when equals 1, it allows correlation between the trend and the cycle (default, 0)
- `opts.mu0` (`opts.Vmu`): prior mean (variance) trend drift (default, 0.1 [(default, 1)])
- `opts.phi0`: prior mean cycle AR1 and AR2 coefficients (default, [0.2 0.2])
- `opts.Vphi`: prior variance AR coefficients (default, diagonal matrix with 1)
- `opts.sigc2` (`opts.sigtau2`): prior mean cycle (trend) variance (default, 0.5) [(default, 1.5)]

The approach uses the MCMC implementation of Grant and Chan (2017). 

Trend and Cycle decompositions III

- **Unobservable component (UC) decomposition 2.** The trend is assumed to have a local linear specification, the cycle is an AR(p), where p selected by the user. Estimation is by maximum likelihood-Kalman filter. The decomposition can be activated using:

$$[yt, yc, out] = uc2_(y, lags, opts);$$

where y is the data, lags is the number of lags in the cycle. The parameters are ordered as: $[\varphi_1, \dots, \varphi_p, \sigma_\epsilon, \sigma_{\eta,1}, \sigma_{\eta,2}]$ (AR, stdev of measurement error, stdev of latent variables); they can be estimated or calibrated by setting the appropriate options and are reported as additional subfields in out. opts is a field structure that allows customization:

- opts.phi is a $lags \times 1$ vector that sets the initial values for the φ 's (default, 0.5^j for $j = 1, \dots, d$).
- opts.sigma is a 3×1 vector that sets the initial values for the σ 's (default, 0.5).
- opts.index_est is a row vector that selects the parameters to be optimized. By default, all parameters are optimized, $opts.index_est=1:lags+3$.
- opts.lb and opts.ub set the lower and upper bounds for the optimization. Both are row array vectors of the same size of opts.index_est.
- opts.max_compute is a scalar selecting the maximization routine. The settings are the same as those for the maximization of the Minnesota hyper-parameters (see manual and later).

Trend and Cycle decompositions IV

- **Beveridge and Nelson (BN) univariate.** The decomposition is activated using:

$$[dc, dt] = \text{BNuniv}(y, \text{lags}, ff, gg, mm);$$

where y is the time series, lags is the number of lags in the estimated AR; ff is an indicator function for the constant ($ff = 0$ no constant); gg is an indicator function for whether the estimated cycle is plotted or not ($gg = 0$ no plot); mm is an indicator function for whether the decomposition is computed using the model-based long run mean or the estimated mean of y_t ($mm = 1$ use the estimated mean; $mm = 0$ use long run mean).

- **Hodrick and Prescott filter.** The decomposition is activated using:

$$dt = \text{Hpfilter}(y, lam); \quad dc = y - dt;$$

where y is the data, lam is a smoothing parameter. The outputs are the estimated trend dt . lam must be inputted by the user (there is no default value). Typical choice are 1,600 for quarterly data, 6.25 for annual data, and 129,000 for monthly data.

- **One-sided Hodrick and Prescott filter.** The decomposition is activated using:

$$[dt, dc] = \text{one_sided_hpfilter_serial}(y, lam, disc);$$

where y is the data, lam is the smoothing parameter, $disc$ the number of unused observations (typically, $disc = 0$). The outputs are the estimated trend dt and the estimated cycle dc . The one-sided HP filter only looks at past and current data to construct trend estimates (rather than past, current and future data).

Trend and Cycle decompositions V

- **Band Pass filter: Baxter and King implementation** It can be activated using:

$$dc = bkfilter(y, bp1, bp2); \quad dt = y - dc;$$

$bp1$ ($bp2$) is the upper (lower) limit of the frequency band over which the squared gain function is 1. No default values for $bp1$ and $bp2$ are set. Typical choices for quarterly data are $bp1 = 8$ and $bp2 = 32$. For monthly (annual) data typical choices are $bp1 = 24$ (2) and $bp2 = 96$ (8).

- **Band Pass filter: Christiano and Fitzgerald implementation** It can be activated using:

$$dc = cffilter(y, bp1, bp2, cf1, cf2, cf3); \quad dt = y - dc;$$

where $cf1$ is an indicator function describing the integration properties of the input ($cf1 = 0$ no unit root); $cf2$ is an indicator function for whether there is a drift in the unit root or a deterministic trend ($cf2 = 0$ no drift); $cf3$ chooses the format of the filter:

- $cf3 = 0$ uses the basic asymmetric filter;
- $cf3 = 1$ uses the symmetric filter;
- $cf3 = 2$ uses a fixed length symmetric filter;
- $cf3 = 3$ uses a truncated, fixed length symmetric filter (the BK choice);
- $cf3 = 4$ uses a trigonometric regression filter.

If $cf1$, $cf2$ and $cf3$ are omitted, $cf1 = 1$, $cf2 = 1$ and $cf3 = 0$ are the defaults.

Trend and Cycle decompositions VI

With the Great Moderation,
the cycles tend to be longer.

- **Wavelet filter.** The decomposition is activated using:

$$[dc1, dt1, dc2, dt2] = \text{wavefilter}(y, gg);$$

where $dc1$ and $dc2$ are the cyclical components computed focusing on 8-32 or 8-64 quarters cycles and $dt1 = y - dc1$, $dt2 = y - dc2$ are the corresponding trend estimates; y is the input series, and gg is an indicator function for whether the cyclical components are plotted (for $gg = 1$ plots are shown).

- **Butterworth filter.** No option set; implementable with MATLAB commands (see manual).

Cuts spectrum horizontally, whereas wavelet vertically.

- **MA filter.** No option set; implementable with MATLAB commands (see manual).

Trend and Cycle decompositions VII

- **Bivariate Beveridge and Nelson and Blanchard and Quah decompositions.** The two decompositions are jointly activated using:

$$[dt1, dc1, dt2, dc2] = \text{BNBQbiv}(y, \text{lags}, ff, \text{rhoyes}, \text{ssts}, gg, mm);$$

where lags is the number of lags in the VAR; ff is an indicator function for the constant (if $ff = 1$ there is a constant); rhoyes is an indicator function for whether the two disturbances are correlated (if $\text{rhoyes} = 1$ shocks are correlated; this option triggers the Cover et al. (2006) decomposition); ssts is an indicator function for whether the other VAR variables are also integrated (if $\text{ssts} = 1$ it is $I(1)$); gg is an indicator function for the plots (if $gg = 1$ the permanent and transitory component of the first variable are plotted); mm is an indicator function for spectral density plots (if $mm = 1$ spectral densities plots appear on the screen). y is a $T \times 2$ matrix; $dt1$, $dc1$ ($dt2$, $dc2$) are the permanent and the transitory components of the first series obtained with BN (with BQ).

Practice I

Sample codes are in ...\\BVAR\\examples\\Trend-Cycle-dating tutorial.

- Dating of turning points: US and Euro area data (example_2_dating.m).
- Linking the cyclical components of output and credit (example_1_cycles.m)

Run them, do the little exercises, and acquaint yourself with the language and the syntax.

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference**
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Classical VAR Estimation and Inference

- Estimation in EMT can be performed with Classical or Bayesian methods. By default, a Jeffrey (flat) prior is assumed, which results in classical estimates.
- Given a prior, draws from the posterior distribution of the parameters are produced using a Gibbs sampler algorithm; see the guide for more details.
- The baseline, all purpose, estimation function is

$$[\text{BVAR}] = \text{bvar_}(y, \text{lags}, \text{options})$$

- Inputs are:
 - y a $(T \times n)$ array of data (rows=time; columns=variables)
 - $\text{lags} > 0$ the number of lags.
 - options is a field structure that specifies various customized options (i.e. priors, horizons, identification schemes). It can be omitted. The default identification scheme is Cholesky.

Outputs

- BVAR a field with many sub-fields; the most important are (check manual for full list).
 - `BVAR.Phi_draws` is a $(n \times \text{lags} + 1) \times n \times K$ array containing K draws for the parameters. `BVAR.Phi_draws(:, :, k)` stacks vertically the AR matrix; last row of `BVAR.Phi_draws(:, :, k)` contains the constant, Φ_0 , if it assumed to be present.
 - `BVAR.Sigma_draws` is a $n \times n \times K$ matrix containing K draws for Σ .
 - `BVAR.e_draws` is a $(T - \text{lags}) \times n \times K$ matrix containing K draws for the innovations.
 - `BVAR.Phi_ols`, `BVAR.Sigma_ols` and `BVAR.e_ols` are the OLS point estimate analogs.
 - `BVAR.ir_draws` is a four dimensional object (i.e. $n \times \text{hor} \times n \times K$ matrix) that collects the impulse response functions with the chosen identification.
 - `BVAR.forecasts.no_shocks` (`with_shocks`) is a three dimensional object (i.e. $\text{fhor} \times n \times K$ matrix) that collects the forecasts assuming zero (non-zero) shocks in the future.
 - `BVAR.logmlike` contains the log marginal likelihood — a measure of fit.
 - `BVAR.InfoCrit`: contains the Akaike, AIC, the Hannan-Quinn HQIC and the Bayes BIC, information criteria.
- The last two options can used to select the lag length: the log marginal likelihood measures one step ahead predictive ability; other criteria measure in-sample fit.
- By default, $K=5000$ (change it with `options.K = number`) and a constant is assumed (change it with `options.noconstant = 1`).

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference**
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Priors

Three types of priors are allowed in EMT:

- Uninformative or Jeffrey priors (default, no further instruction needed).

```
[BVAR] = bvar_(y, lags, options)
```

- Minnesota prior with default hyper-parameter values.

```
options.priors.name = 'Minnesota';
```

```
[BVAR] = bvar_(y, lags, options)
```

- Multivariate-Normal Inverse-Wishart Conjugate with default values.

```
options.priors.name = 'Conjugate';
```

```
[BVAR] = bvar_(y, lags, options)
```

- If no other options are specified, default parameters are used.

Minnesota prior

- 5 scalar hyper-parameters control the shrinkage of the Minnesota prior:
 - τ `options.minn_prior_tau`: overall tightness (default 3). The larger is τ , the tighter is prior.
 - d `options.minn_prior_decay`: tightness on the lags greater than one (default 0.5). The larger is d , the faster is the lag decay.
 - λ `options.minn_prior_lambda`: the sum-of-coefficient prior tightness (default 5).
 - μ `options.minn_prior_mu`: co-persistence prior tightness (default 2).
 - ω `options.minn_prior_omega`: covariance matrix tightness (default 2).
- Three ways to activate the Minnesota prior by defining the appropriate option command:
 - 1 default values:
`options.priors.name = 'Minnesota'`
 - 2 cherry-picking hyper-parameter values:
`options.minn_prior_tau=5;`
`options.minn_prior_lambda=0.01;`
 - 3 optimized hyper-parameter values:
`options.max_minn_hyper = 1; % default for maximization`

Minnesota prior with optimal hyper-parameters

- Various options can be set for the maximization step:
 - `options.index_est` is a row vector selecting the parameters to be optimized (default, `options.index_est=1:5`)
 - `options.lb (.ub)` sets the lower (upper) bound for the optimization. It has the same size of `options.index_est`.
 - `options.max_compute` is a scalar selecting the maximization routine to be employed:
 - `options.max_compute = 1` uses the MATLAB `fminunc.m`
 - `options.max_compute = 2` uses the MATLAB `fmincon.m`
 - `options.max_compute = 3` uses the Chris Sims's `cminwel.m` (default)
 - `options.max_compute = 7` uses the MATLAB `fminsearch.m`
- Once the maximum is found, the posterior distribution is computed using optimal hyper-parameter values. If optimization is unsuccessful, the posterior distribution is computed with default values (**and a warning is issued**).
- Tip: it is a good idea to max one parameter at time, starting the optimization from the values obtained in the previous step (`bvar_max_hyper` can be used to shorten the computational time. It does hyper-parameter optimization, rather than full Gibbs sampling computations).

```
[hyperp,~,~] = bvar_max_hyper(x0,y,lags,options);
```

Conjugate MN-IW

- A multivariate Normal-Inverse Wishart conjugate prior can be activated using:

```
options.priors.name = 'Conjugate'
```

Default values: $\Phi \sim N(0, 10 \, I_{np+1})$ and $\Sigma \sim IW(I_n, n + 1)$

Useful options:

- `options.priors.Phi.mean` is a $(n \times \text{lags} + 1) \times n$ matrix containing the prior means for the autoregressive parameters.
- `options.priors.Phi.cov` is a $(n \times \text{lags} + 1) \times (n \times \text{lags} + 1)$ matrix containing the prior covariance for the autoregressive parameters.
- `options.priors.Sigma.scale` is a $(n \times n)$ matrix containing the prior scale of the covariance of the residuals (use it to adjust default scale matrix if variables have different units).
- `options.priors.Sigma.df` is a scalar defining the prior degrees of freedom.

VARX

- Can estimate Classical or Bayesian VAR with exogenous variables (X). To allow exogenous variables need to specify it as an option

```
options.controls =x
```

where x is a $(T \times q)$ matrix containing the exogenous variables

q can be lag or contemporaneous

- The first dimension of x is time and must coincide with the time dimension of y (or with the sum of the time dimension of y and the out-of-sample forecast horizons.)
- Lags of exogenous controls can be used and specified as additional columns in x .
- A VARX model can be estimated with Jeffrey priors (classical) or conjugate priors (Bayesian); Minnesota priors are not currently supported.
- To launch the estimation use the standard command:

```
BVAR = bvar_(y,lags,options);
```

- Draws with the responses of the endogenous variables to shocks to the exogenous variables are stored in BVAR. `irx_draws` (a four dimensional matrix). If there is more than one exogenous variable, shocks are computed using a Cholesky decomposition.

Regularization methods

- When `bvar` is used for the estimation of a large system, it allows the use of regularization methods.

Elasticnet typically preferred, which is why it was included...

- Regularization methods can be initialized with the option command: e.g.
`options.Ridge.est=1,options.Lasso.est=1, options.ElasticNet.est=1.`

quadratic penalty

absolute constraint

both quadratic & absolute

- Parameters can be set in option command, e.g. `options.Ridge.lambda=0.02,`
`options.ElasticNet.alpha=0.05.`

the rate of the penalty (default)

- Estimation uses the standard command `bvar1= bvar_(y, lags, options).`
- If any of the standard priors is used, it produces K draws for the statistics of interest, e.g.
`options.K = 1000.`

All biased, squared error

Network analysis I

connectedness = how much
of a shock in average gets
transmitted in the system

- Statistics computed using the variance decomposition at horizon `option.nethor=H`.
- Default is Pesaran-Shin GIR decomposition: `options.connectedness = 1`.
- Other identification methods are possible, e.g.

```
options.connectedness = 2;      % activate alternative identification schemes
```

```
options.signs{1} = 'y(a,b,c)>0';    % activates sign restrictions
```

`bvar1` contains:

- a measure of overall connectedness: `bvar1.Ridge.Connectedness.Index`;
- a measure of spillovers: `bvar1.Lasso.Connectedness.FromUnitToAll`;

how the capital of the bank changes
- a measure of vulnerability: `bvar1.ElasticNet.Connectedness.FromAllToUnit`.

if `options.connectedness = 1`, outputs
all three things

Practice II

Sample codes are in ...\\BVAR\\examples\\BVAR tutorial.

- Calculation of the optimal lag length; Cholesky responses and variance decomposition of monetary policy shock (example_1_classical.m)
- Estimation with Minnesota prior with hyper-parameter optimization (example_2_minn.m)
- Effects of US interest rate shocks on UK variables (example_6_VARX.m)
- Measuring network effects in cryptocurrency markets (example_11_connetdness.m)

start from the optimal variable, before putting them all in, otherwise will not work in all systems (..?..)

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)**
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

IRF

- For each draw of the posterior distribution, EMT computes the IRF for a given identification scheme.
- IRFs are stored in a four dimensional arrays of dimension $(n \times hor \times n \times K)$ (variable, horizon, shock, draw).
- Flat, Minnesota, Conjugate MN-IW priors can be used.
- Default value for the horizon is 24; it can be changed with `options.hor = hor; .`
- By default Cholesky decomposition is used: `BVAR.ir_draws` is the default matrix containing IRF (no further command needed).
- Other identification schemes need to be specified by the user in the `options`, before launching `bvar`.

Plotting IRFs

- The plotting command for IRF is :

```
plot_irfs(irfs_to_plot, optnsplt)
```

- `irfs_to_plot` is a four dimensional array; e.g. `irfs_to_plot = BVAR.ir_draws` or `irfs_to_plot = BVAR.irproxy_draws(:, :, 1, :)`;
- `optnsplt` is optional and defines plotting options. The most relevant are:
 - `optnsplt.varnames (.shocksnames)` is a cell string with variable (shock) names.
 - `optnsplt.conf_sig (.conf_sig_2)` is a number between 0 and 1 indicating the size of (second) highest posterior density (HPD) set to be plotted; the default is 0.68 for the first (no default for the second).
 - `optnsplt.saveas_strng (.saveas_dir)` is a string array with name of the directory where to save the plot(s).
 - `optnsplt.add_irfs` allows to add other IRFs to the plot when one wants to compare responses.
 - `optnsplt.nplots` is a $n \times m$ array indicating the structure of the subplots.

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions**
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Long run restrictions

- Activated using:

```
options.long_run_irf = 1;
```

- Need the first variable in the y vector to be specified in log difference
- By default, the first disturbance has long run effects on the first variable.
- Basic setup follows Gali (1999).
- If more than one shock affects the variables in the long run, as in Fisher (2006), use more than one variable in first difference.
- IRFs are stored in the matrix `BVAR.irlr_draws`

Sign and magnitude restrictions

- Activated using:

```
options.signs{1} = 'y(a,b,c)>0'
```

where the array 'y' refers to the IRF, 'a' to the variable, 'b' to the horizon, 'c' to the shock.

- The syntax means that the shock c has a positive impact on variable a at horizon b.
- The syntax is flexible and allows several variations (e.g. lower bound on elasticity or threshold on the cumulative impact)

horizon = 1
(mid)

```
options.signs{2} = 'y(a_1,1,c)/y(a_2,1,c) >m'
```

```
options.signs{3} = 'max(cumsum(y(a,:,c),2))<M'
```

- Implementation follows: Rubio-Ramirez, Waggoner and Zha (2010).
- IRFs are stored in the matrix `BVAR.irsign_draws`.

Narrative restrictions

historical decomposition uses averages, which may not be representative with the standard errors

- Activated using:

```
options.narrative{1} = 'v(tau,n)>0'
```

where the array '*v*' refers to the structural innovation, '*tau*' to time, and '*n*' to the shock.

- The syntax means that shock *n* is positive on the time *tau* periods.
- Again the syntax is flexible and many variations are allowed:

```
options.narrative{1} = 'sum(v(tau_0:tau_1),n)>0'
```

Here the sum of the shock *n* between periods *tau_0* and *tau_1* is required to be positive.

- Implementation follows Ben Zeev (2018) and Antolin-Diaz and Rubio-Ramirez (2018)
- IRFs are stored in the matrix `BVAR.irnarrsign_draws`.

IV approach

- External instrumental or proxy variable estimation is activated using :

```
options.proxy =instrument;
```

- Restriction 1: Instrument vector cannot be longer than $(T - lags)$
- Restriction 2: Instrument vector ends when the VAR ends (default). When this is not the case, use the option

```
options.proxy_end =periods;
```

to line up instruments and data. `periods` is the number of periods separating the last observation of the instrument and the last observation of the VAR innovations.

- Multiple proxy variables are allowed to identify one structural shocks, i.e. `instrument` can be a vector. If one is interested in identifying, say, two disturbances, she need to repeat the IV exercise twice, separately for each disturbance (currently, no option to identify multiple structural shocks with multiple proxy variables).
- Implementation follows Miranda-Agrippino and Ricco (2021)
- IRFs are stored in the matrix `BVAR.irproxy_draws`
- By convention, the structural shock of interest is the **first**; i.e.
`BVAR.irproxy_draws(:, :, 1, :);`

Mixed identification restrictions

A mix of zero (short and long run) and sign restrictions can be used. This option is activated using a combination of previous options:

- `options.zero_signs{1} = 'y(j,k)=+1'`
Here shock k has a positive effect on the j -th variable **on impact**
- `options.zero_signs{2} = 'ys(j,k_1)=0'`
Here shock k_1 has a zero impact effect on the j -th variable.
- `options.zero_signs{3} = 'yl(j,k_2)=0'`
Here shock k_2 has a zero long run effect on the j -th variable.
- **Pay attention: short and long run restrictions are distinguished by the addition of r or s to the y vector.**
- Implementation follows Arias, Rubio-Ramirez, Waggoner and Zha (2018) and Binning (2013).
- IRFs are stored in the matrix `BVAR.irzerosign_draws`

Heteroskedasticity identification

- Relative variances of structural shocks change across 'regimes'. This could be a source of identification of the coefficients of a structural VAR.
- Implementation follow Sims (2020)
- Specify a $T \times 1$ vector of the same size y , containing zero and ones, where zero (one) indicate the first (second) regime (Right now more than two regimes not allowed).
- Heteroskedasticity based responses are obtained activating the options:

```
options.heterosked_regimes = Regimes;
```

where `Regimes` is $T \times 1$ vector with zeros and ones.

- Estimated responses and the rotations are stored in, respectively

```
BVAR.irheterosked_draws      BVAR.Omegah;
```

`BVAR.irheterosked_draws` is a four dimensional object (i.e. a $n \times \text{hor} \times n \times K$ matrix): the first dimension corresponds to the variables, the second to the horizon, the third to the disturbances, and the fourth to the responses produced by a posterior draw. `BVAR.Omegah` is a three dimensional object (i.e. a $n \times n \times K$ matrix): the first dimension corresponds to the variables, the second to the shocks of interest, and the third to a posterior draw.

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions**
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Other summary functions

- The forecast error variance decomposition is activated using:

```
FEVD = fevd(hor,Phi,Sigma,Omega);
```

FEVD is a $n \times n$ array; the (i,j) element the share of variance of variable i explained by shock j at horizon hor . $Omega$ is the rotation matrix used (omit, if Cholesky decomposition is employed).

- The historical decomposition is activated using:

```
[yDecomp,v] = histdecomp(BVAR,opts);
```

where $yDecomp$ is a $T \times n \times n+1$ array (time, variable, shock and initial condition $\bar{\Psi}_t$); v is the $T \times n$ array of structural innovations; $BVAR$ is the output of the `bvar.m` function; `opts.Omega` declares a different rotation/identification (omit, if Cholesky decomposition is employed).

- To plot the historical decomposition use:

Bug: clears all figures, but they can be found from the subfolder.

```
plot_shcks_dcmp(yDecomp, BVAR, optnsplt)
```

`optnsplt` is optional and controls various plotting options (see manual).

Additional identifications: Maximize FEVD at horizon h

- Given a draw (Phi, Sigma) of the reduced form VAR parameters, the command:

```
Qbar = max_fevd(i, hor, j, Phi, Sigma)
```

can be used to find the orthonormal rotation maximizing the forecast error variance decomposition (FEVD) of variable i explained by shock j at horizon hor.

- Given the selected Qbar, IRFs can be computed using

```
y_fevd = iresponse(Phi, Sigma, hor, Qbar)
```

- Example: assume the shock of interest is the first and we want it to explain the largest portion of the variance of the first variable four periods ahead:

```
j = 1; i=1; h=4; % shock, variable, period
for k = 1 : BVAR.ndraws % iterate on posterior draws
    Phi = BVAR.Phi_draws(:, :, k);
    Sigma = BVAR.Sigma_draws(:, :, k);
    Qbar = max_fevd(i, h, j, Phi, Sigma);
    [ir] = iresponse(Phi, Sigma, hor, Qbar); %responses with a posterior draw
    BVAR.irQ_draws(:, :, :, k) = ir; % store the IRFs
end
```

- Another useful function: `Q = generateQ(n)`

It is a rotation matrix generator which can be used in

`y = iresponse(Phi, Sigma, hor, Q)` to produce impulse responses for a generic rotation Q.

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations**
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Time variations

- No direct command to estimate parametric time varying BVARs. Can however do non-parameteric (fixed rolling window) estimation to check for time variations in the statistics of interest. See example 16 in the manual (page 49), or `example_3_IRF.m` in the tutorials.
- No direct command to control for time varying volatility. One can however estimate the model correcting for abnormal volatility movements (like those produced by COVID19) by playing around with the available options. See example 17 in the manual (page 52) or `example_10_VAR_heterosked.m` in the tutorials.

Practice III

Sample codes are in ...\\BVAR\\examples\\BVAR tutorial.

- Calculation IRFs with various identification schemes. Computation of variance and historical decompositions (example_3_irfs.m).
- Estimation of a TVC BVAR (example_3_irfs.m).
- Estimation of a BVAR with volatility changes (example_10_VAR_heterosked.m) .

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM**
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Larger scale models

- For large scale VARs, the best option is a Minnesota prior with a value of τ (overall tightness) larger than the default, to a priori set many coefficients close to zero.
- The factor structure of Canova and Ciccarelli (2009)-(2013) is not implemented here. See the BEAR toolbox, if you want to use it.
- To compress a vector of time series into static principal components, use

```
[E, W, Lambda, EigenVal, STD] = pc_T(y_1, n_w, transf);
```

`y_1` contains the data to be compressed, a $T \times n$ vector, `n_w` is a scalar, indicating the number of factors to be extracted, and `transf` indicates the data transformation: 0 means no transformation, 1 means demean, and 2 means demean and standardize. The outputs `E`, `W`, `Lambda`, `EigenVal` are, respectively, the idiosyncratic errors, the principal components, the loadings and the eigenvalues. When `transf=2`, `STD` is the $n_2 \times 1$ vector containing the standard deviation of `y_2`; otherwise, it is a vector of ones.

- Alternatively, to compress time series (with similar units and similar standard deviations) one can compute a cross sectional mean at each `t`:

```
y_2=mean(y1,2,'omitnan') or y_2=trimmean(y_1, pct,'round',2)
```

where `pct` is the percentage of observations eliminated.

FAVAR models

- Once y_1 or y_2 are computed, they can be added to the VAR as factors using

```
FAVAR = bvar([W y_3], lags)           FAVAR = bvar([y_2 y_3], lags) ;
```

where W are the principal components, y_2 are cross sectional means and y_3 variable excluded from y_1 .

- Any prior specification and any identification restriction could be used.
- When demeaned standardized principal components are used, statistics for the original variables can be computed using two commands jointly:

```
ReScale = rescaleFAVAR(STD,Lambda,n_1) ;
```

where $ReScale$ is a $(n_2+n_1) \times (n_w + n_1)$ matrix. By default, factors loadings are ordered first. This rescales the FAVAR coefficient estimates.

Then for any statistics, take the output of FAVAR and transform it with $ReScale$. For example., assuming interest is in the mean forecast of y_2 , type

```
mean_forecast_y = mean(fabvar.forecasts.no_shocks,3) * ReScale';
```

Principal components

- Principal Components (PC) are linear combinations of observables maximizing the explained portion of their variance. For any $T \times n$ matrix of iid random variables Y with covariance Σ_y , we find the linear combination $Y\delta$, where δ is a $n \times m$ vector, such that the variance of $Y\delta$ (i.e. $E(\delta' Y' Y \delta) = E(\delta' \Sigma_y \delta)$) is maximized subject to $\delta' \delta = I$.
- The function `pc_T.m` allows to extract the first principal components:

```
[e, f, Lambda, EigenVal, STD] = pc_T(y, m, transf);
```

Typically, y is assumed to be stationary and standardized; if not, use the `transf` option to transform it within the script.

- There is no agreement on how to select the number `m` of principal components. Typically, since eigenvalues are reported in decreasing order, one selects a desired fraction of total variance to be explained and include as many PC as it is necessary to reach that level (recall that the fraction of variance explained by first m factors is $\frac{\sum_{i=1}^m \lambda_i}{\text{trace}(\Sigma)}$).
- Scree plots or the Bai-Ng (2002) test can also be used.

Static factor models

- A model with $m < n$ static factors can be represented as:

$$y_t = \Lambda f_t + e_t$$

where y_t is a $n \times 1$ vector, f_t is a $m \times 1$ vector of common factors with zero mean and covariance matrix Σ_f ; e_t is a $n \times 1$ random vector of idiosyncratic errors with zero mean and diagonal covariance matrix Σ_e ; Λ is the $n \times m$ matrix of factor loadings.

- The factors f_t could be observable or latent. If they are observables the model is simply a multivariate regression. When they are unobservable, they can be estimated with principal component techniques or Bayesian methods.
- Latent factors f_t and loadings Λ are not separately identified since $\Lambda f_t = \Lambda C C^{-1} f_t = \Lambda^* f_t^*$. Since the number of free parameters in the data is $n(n+1)/2$ and the number of estimated parameters is $nm + m(m+1)/2 + n$, one scheme that generates (over-)identification is to assume $\Sigma_f = I$. Alternatively, one can assume that the $m \times m$ sub-matrix of Λ is normalized to have ones on the main diagonal and zeros on the upper triangular part. Both schemes can be employed for the static factor model; only the latter for the dynamic factor model.
- In EMT the static factor model is a special case of the dynamic factor model. Thus, there is no special command to estimate static factor models.

Dynamic factor models I: General command

- A model with m dynamic factors, with p lags each factor, can be represented as:

$$y_t = \Lambda f_t + e_t$$

$$f_t = \Phi_1 f_{t-1} + \dots + \Phi_p f_{t-p} + u_t$$

where Φ_j , $j = 1, \dots, p$ are $m \times m$ matrices; u_t is an iid normally distributed, zero mean, random vector with covariance matrix Σ_u .

- EMT allows estimation of $\vartheta = [\Lambda, \Sigma_e, \Phi_1, \dots, \Phi_p, \Sigma_u, \{f_t\}_{t=1}^T]$ and inference, conditional on the estimates of ϑ .
- The baseline estimation function is

```
BDFM = bdfm_(y, lags, m, options);
```

y is the data, an array of size $T \times n$; $lags$ is the number of lags of the dynamic factors; if $lags = 0$ a static factor model is estimated; m ; $options$, specifies the options. The data is assumed to be demeaned, unless otherwise specified (use `options.do_the_demean=0`).

- When the f_t and e_t are normal, the model can be estimated with the Gibbs sampler.
- The default priors are: $p(\Lambda)$ is normal with zero mean and with a diagonal covariance matrix of 10; $p(\sigma_{e_i}^2)$ is inverted Gamma prior with unitary scale and four degrees of freedom; $p(\Phi_1, \dots, \Phi_p)$ is normal with zero mean (except for the first own lag, which is set at 0.2) and a diagonal covariance matrix of 10; $p(\Sigma_u)$ is an inverse Wishart with unitary diagonal scale and $m+1$ degrees of freedom. These default values can be changed using previously mentioned options.

Dynamic factor models II: Prior options

- Φ : `options.priors.F.Phi.mean` is a $(m \times \text{lags}) \times m$ matrix containing the prior means for the factors' autoregressive parameters.
- Φ : `options.priors.F.Phi.cov` is a $(m \times \text{lags}) \times (m \times \text{lags})$ matrix containing the prior covariance for the factors' autoregressive parameters.
- Σ_u : `options.priors.F.Sigma.scale` is a $(m \times m)$ matrix containing the prior scale of the covariance of the factor innovations.
- Σ_u : `options.priors.F.Sigma.df` is a scalar defining the prior degrees of freedom for the factor innovations.
- (Φ, Σ_u) : `options.priors.name = 'Jeffrey'`; activate a Jeffrey prior on Φ, Σ_u .
- Λ : `options.priors.F.Lambda.mean` is a scalar defining the prior mean for the factor loadings.
- Λ : `options.priors.F.Lambda.cov` is a scalar defining the prior variance for the factor loadings.
- Σ_e : `options.priors.G.Sigma.scale` is a scalar defining the prior scale of the variance of the error terms.
- Σ_e : `options.priors.G.Sigma.df` is a scalar defining the prior degrees of freedom for the variance of the measurement error terms.

One can set these options one at the time or jointly.

Dynamic factor models III: Content of BDFM

BDFM is a structure with several fields and sub-fields containing the estimation results:

- `BDFM.Phi_draws` is a $(m \times \text{lags}) \times m \times K$ matrix containing K draws for the factors autoregressive parameters, Φ_j . For each k , the first $m \times m$ submatrix of `BDFM.Phi_draws(:, :, k)` contains the autoregressive matrix of order one; subsequent submatrices higher order lags.
- `BDFM.Sigma_draws` is a $m \times m \times K$ matrix containing K draws for the covariance matrix of the factor innovations, Σ_u .
- `BDFM.lambda_draws` is a $n \times m \times K$ matrix containing K draws for the factor loading, Λ .
- `BDFM.sigma_draws` is a $n \times K$ array containing K draws for the variance of the error terms, Σ_e .
- `BDFM.f_draws` is a $T \times m \times K$ matrix containing K draws for the factors, obtained with the Carter and Kohn backward smoothing algorithm.
- `BDFM.f_filt_draws` is a $T \times m \times K$ matrix containing K draws for the factors, obtained with the Kalman filter forward recursions.

Dynamic factor models IV: Content of BDFM

- `BDFM.ef_draws` is a $T \times m \times K$ matrix containing K draws for the factor innovations.
- `BDFM.eg_draws` is a $T \times n \times K$ matrix containing K draws for the measurement error terms.
- `BDFM.ir_draws` is a four dimensional object (i.e. $n \times \text{hor} \times m \times K$ matrix) that collects the impulse responses obtained with a recursive identification on the innovations of the factors. The first dimension corresponds to the variables, the second to the horizons, the third to the disturbances, and the fourth to the responses obtained with particular draw from the posterior distribution of the factor parameters. Alternative identification schemes are possible, if declared prior to the BDFM command.
- `BDFM.forecasts.no_shocks` (`BDFM.forecasts.with_shocks`) is a three dimensional object (i.e. $\text{fhor} \times n \times K$ matrix) that collects the forecasts assuming zero (non-zero) shocks in the factors in the future. The first dimension is the horizon, the second to the variable, and the third to the draw from the posterior distribution of the parameters.

The default value for K is 20,000. It can be changed using `options.K`, prior to calling the `bdfm_` routine.

Dynamic factor models V: Tips

- The number of dynamic factors can be selected using Bai and Ng (2007) approach, i.e. estimate first a VAR in r static factors, where the factors are obtained by PC on a large panel of data, then compute the eigenvalues of the residual covariance (or correlation) matrix and test the rank of the matrix.
- Estimation of factor models may be time consuming; difficult to give general suggestions on how to improve the estimation speed, as the solution tends to be case and data specific. Below are a few tips that might be useful.

- If $T = 200, N \approx 50$, use the default estimation options, keeping in mind that the larger is N , the slower the codes become.
- $T > 1,000$ and $N < 10$, estimates the f_t with the steady state Kalman Gain in the recursion of the Kalman filter. Using the steady state Kalman Gain avoids inverting the prediction error covariance matrix in constructing the updated Kalman Gain (with five years of daily data this avoids $5 \times 365 \times 20,000$ matrix inversions) as the former is computed outside of the time loop. The steady state Kalman Gain can be activated with the option:

```
options.use_the_ss_kalman_gain = 1;
```

- If $N > 100$ (regardless of T), the Law of Large Number applies and one can treat cross-section averages as if they were observed data. One would then estimate a VAR on the cross section averages and use OLS regressions in the measurement equation to recover the dynamics of the original variables.

Dynamic factor models VI: IRFs

- Dynamic factor models are reduced form models. To study the transmission of interesting disturbances, we need to impose some structure.
- Structural disturbances ν_t can be mapped into the u_t , by means of orthonormal rotation matrices, as in VAR models. Thus, all identification schemes for VARs can be used.
 - Choleski identification (default). The responses are in the `BDFM.ir_draws` matrix.
 - Long-run identification. Activated with the option: `options.long_run_irf = 1;` Responses are stored in the `BDFM.irlr_draws` matrix.
 - Sign and magnitude restrictions identification. Activated with the option: `options.signs{1} = 'y(a,b,c)>0'` Responses are stored in the `BDFM.irsign_draws`.
 - Narrative restrictions identification. Activated with the option: `options.narrative{1} = 'v(tau,n)>0'` Responses are stored in the `BDFM.irnarrsign_draws` matrix.
 - External instrumental or proxy variable identification. Activated with the option: `options.proxy = instrument;` Responses are stored in `BDFM.irproxy_draws`.
- Response matrices have four dimensions (i.e. $n \times hor \times m \times K$ matrix). The first corresponds to the variables, the second to the horizons, the third to the disturbances, and the fourth to the responses obtained with particular draw from the posterior distribution of the parameters.
- Numerical approximations to the distribution of the impulse response functions can be computed and visualized using the function `plot_irfs_m`.

Practice IV

Sample codes are in `...\BVAR\examples\BVAR tutorial`.

- Estimation and IRFs in a PCVAR (`example_5_favar.m`) .
- Estimation of a dynamic factor model and computation of IRFs (`example_12_bdfm.m`)

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs**
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Panels of VARs

No specific command to run panels of VARs but many options are allowed with available technology

- Exact pooling of cross sectional data with fixed effects (see example 22, page 69 of the manual).
- Average time series estimator along the lines of Pesaran and Smith (1996) (see example 21, page 66 of the manual).
- Partial (Bayesian) pooling with a data driven (or used supplied) prior mean and variance (see example 23, page 71 of the manual).
- Underlying the examples is the common implicit assumption is that units are 'islands', i.e.
 - o contemporaneous or lagged spillovers across units are allowed.
- Can use VARX technology to account for common (areawide) fluctuations or can scale unit specific variables by aggregate variables.

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs**
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix

Mixed Frequency VAR

Same technology can be used when combining weekly data with monthly, for example (of course we need to be careful that weekly data matches the monthly).

- Mixed Frequency VAR (MF-VAR) can be used to nowcasts, backcast or to retrieve variables that have an arbitrary pattern of missing observations.
- Examples: Construct monthly GDP, weekly Unemployment.
- Estimation is performed with the Gibbs sampler:
 1. Generate a draw from the posterior distribution of the reduced form VAR parameters, conditional on observables variables and the states.
 2. With the draw run the Kalman filter and the Kalman smoother and get estimate the unobserved states (monthly GDP).
 3. Repeat 1. and 2 for each draw.
- It allows flat, Minnesota (except maximization), Conjugate MN-IW priors. It allows all identification schemes for IRFs. It allows all forecasting options.
- MF-VAR estimation is automatically triggered in the `bvar` function whenever there are NaN in the data, `y`. (A warning is printed on the screen).

triggers only when there's a "hole" in the middle

Example: VAR with Monthly-Quarterly data

- Suppose the time unit is a month.
- Let $y = [y^m, y^q]$ where the y^q variable is observed only in the last month of the quarter q , elsewhere is NaN. Define the unobserved states as x_t .
- Use a state space representation where the transition equation is a monthly VAR for the x_t and measurement equation maps x_t into y_t .
 - For **monthly** variables, the observation equation is: $y_t^m = x_{m,t}$
 - For **Stock** quarterly variables, the observation equation is : $y_t^q = x_{q,t}$. [no instruction is needed in this case]
 - For **Flow** quarterly variables (e.g. GDP), observation equation is:
 $y_t^q = \frac{1}{3}(x_{q,t} + x_{q,t-1} + x_{q,t-2})$. In this case one needs to specify the option:

`options.mf_varindex` = `num`

where `num` is a scalar with the position of the q flow variable (column number in y)

- Additional outputs produced by mixed frequency estimation:
`BVAR.yfill` (smoothed), `BVAR.yfilt` (filtered) $T \times n \times K$ arrays containing smoothed (filtered) estimates of the latent variables.

Practice V

Sample codes are in ...\\BVAR\\examples\\BVAR tutorial.

- Calculation of average IRFs in a panel of VARs (example_8_panels.m).
- Estimation and inference in a mixed frequency BVAR (example_4_mfvar.m).

Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts**
- 14 Local projections and Direct forecasts
- 15 Appendix

Forecasts

Forecasts are automatically computed with BVAR. They are stored as follows:

- `BVAR.forecasts.no_shocks` is a $(nfor \times n \times K)$ matrix (time, variable, draw). All future shocks are zeros (this are unconditional forecasts).
- `BVAR.forecasts.with_shocks` is a $(nfor \times n \times K)$ matrix (time, variable, draw). All future shocks are drawn randomly from the posterior distribution.
- `BVAR.forecasts.conditional` is a $(nfor \times n \times K)$ matrix containing the conditional forecasts. Need to set a bunch of options:
 - `options.endo_index` is a row array containing the index of the variable constrained to a specified path.
 - `options.endo_path` is a matrix containing the path for each variable (rows: horizon, columns: variables). `size(options.endo_path,1) = options.fhor`.
 - `options.exo_index` [optional] specifies the shocks of the VAR used to generate the assumed paths of the endogenous variables. `exo_index` could be one or more shocks. If no structural identification is performed, by default, the option uses a Cholesky factorization.
- `BVAR.forecasts.EPS` contains the shocks used to generate the conditional forecasts.
- Can use `options.fhor` to change the forecasting horizon

Plotting the Forecasts

- The command to plot forecast is:

```
plot_frst_(frst_to_plot,y,T,optnsplt);
```

- `frst_to_plot` is a three dimensional array (time,variable,draw) containing the forecast; e.g. `frst_to_plot = BVAR.forecasts.no_shocks`.
- `y` is the $(T \times n)$ array of data.
- `T` is the $(T \times 1)$ array with the in-sample time span; e.g. for quarterly data, `T= 1990 : 0.25 : 2010.25` if the sample is 1990Q1-2010Q2.
- `optnsplt` can be omitted and defines a number of options. Many are similar to those in `plot_irf_`. A new useful one is:
- `options.order_transform` is a $1 \times n$ array with values
 - `=0` → no transformation performed.
 - `=1` → period-by-period change
 - `=12` → 12 periods change (monthly data) multiplied by 100, i.e. $100(y_{t+12} - y_t)$.
 - `=4` → 4 periods change (quarterly data) multiplied by 100, i.e. $100(y_{t+3} - y_t)$.
 - `=100` → one period change (annual data) multiplied by 100.

More sophisticated than just
dummying stuff out.

Forecasting in times of pandemics

- Are VAR (time series) models useful to forecast during the pandemic?
- Schorfheide and Song (2020): MF-VAR can be used to generate real-time macroeconomic forecasts during the COVID-19 pandemic. Quite pessimistic outlook for the US: long-lasting recession.
- Primiceri and Lenza (2020): **Heteroskedasticity adjusted VAR**. Scale down the observables during the peak of the COVID-19 pandemic. Seems a sensible idea.

$$y_t = \mathbf{x}_t' \Phi + \frac{1}{w_t} u_t$$

w_t is a vector of weights. Weights can be picked by the investigator. Implementation

- `options.heterosked_weights` is the $(T - \text{lags} \times 1)$ array of weights; e.g. if we are interested in scaling sample 2000m2 to 2020m7, use:

```
options.heterosked_weights = [1 .... 1 20 30 30 20 20 10 1 .....1];
```

..., 1/20, 1/30,

- Weights can also be optimized (as Minnesota hyper-parameters) by defining
- ```
options.objective_function = 'bvar_opt_heterosked';
options.tstar = find(time==2020) + 2; %picks March 2020
[postmode,logmlike,HH] = bvar_max_hyper(hyperpara,y,lags,options);
```

# Practice VI

Sample codes are in ...\\BVAR\\examples\\BVAR tutorial.

- Forecast produced in a BVAR (example\_9\_prediction.m ).
- Forecast in time of a pandemic (example\_10\_VAR\_heterosked.m).

# Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts**
- 15 Appendix

# Classical Local projections I

y can be multiple variable (vector)

- The baseline estimation function for local projections is:

`[DM] = directmethods(y, lags, options);`

The first input,  $y$ , is the data (no missing values are allowed); the second,  $lags$ , is the number of lags of  $y$  used in the projection; the third,  $options$ , specifies the options; it can be omitted, if default options are used. Available options are:

- `options.hor` is a scalar indicating the horizon of the responses (default 24).
- `options.conf_sig` is a number between 0 and 1, with the size of the confidence interval (default 0.9).
- `options.controls` is a  $(T \times n_c)$  array of  $n_c$  exogenous controls and the first dimension of controls must coincide with the first dimension of  $y$ .
- `options.proxy` is a  $(T \times n_p)$  array containing  $n_p$  variables that proxy for the shock of interest. The first dimension must coincide with the first dimension of the endogenous variables,  $y$ .
- `options.robust_se_` a scalar indicating the type of standard errors to be computed; if 0, no adjustment is made; if 1 standard errors are HAC adjusted as in Hamilton (1994), Ch 10, pag 282, (default); if 5, the MATLAB `hac.m` function is used (the Matlab Econ Toolbox required).
- `options.Q` is a  $(n \times n)$  orthonormal matrix,  $Q'Q = QQ' = I$ , that can be used to rotate the Cholesky impact matrix. EMT assumes that  $\Omega = chol(\Sigma_{(0)})Q$ , and the default is  $Q = I$ .

for rotating  
the impulse  
responses



# Classical Local projections II

The output of the function, `DM`, is a structure with several fields and sub-fields. The two most important ones are:

- `DM.ir_lp` is a  $(n \times \text{hor} \times n \times 3)$  array containing the impulse response function computed with a recursive identification scheme. The first dimension corresponds the endogenous variable, the second to the horizon, the third to the shock. The first and third elements in the fourth dimension are to the upper and lower limits of the confidence interval, whereas the second dimension is the mean.
- `DM.irproxy_lp` a  $(n \times \text{hor} \times 1 \times 3)$  array containing the impulse response function computed with a proxy shock. The dimensions are the same as above.

Forecasts are also stored in `DM`; in particular:

`DM.forecasts;`

is a  $(n \times \text{hor} \times 3)$  array containing the forecasts. The first dimension corresponds the endogenous variable, the second to the horizon. The first and third elements in the third dimension correspond to the upper and lower limits of the confidence interval, whereas the second element of the third dimension is the mean.

# Bayesian Local projections I

Bayesian local projections with a Normal-Inverted Wishart prior can be acticated with:

- `options.priors.name = 'Conjugate'`

remember: variables should have the same scale (unit), i.e. millions and percentages won't mix

With this setting, the prior for AR parameters has zero mean zero and diagonal covariance matrix of 10 and  $\tau_h = 1$ ; the prior for the covariance matrix of the residual is inverse Wishart with a unitary diagonal matrix as scale and  $n+1$  degrees of freedom.

Customization of prior parameters can be done as follows:

- `options.priors.Phi.mean` is a  $(n \times \text{lags} + 1) \times n$  matrix containing the prior means for the autoregressive parameters.
- `options.priors.Phi.cov` is a  $(n \times \text{lags} + 1) \times (n \times \text{lags} + 1)$  matrix containing the prior covariance for the autoregressive parameters.
- `options.priors.Sigma.scale` is a  $(n \times n)$  matrix containing the prior scale for the covariance of the residuals.
- `options.priors.Sigma.df` is a scalar defining the prior degrees of freedom.
- `options.priors.tau` is a  $(\text{hor} \times 1)$  vector controlling shrinkage at various horizons; the default is  $\tau = 1$ . The larger is  $\tau$ , the larger is the shrinkage around the prior mean.

# Bayesian Local projections II

`options.priors.tau` can also be chosen to maximize the marginal data density using:

```
options.priors.max_tau = 1;
bdm_opt = directmethods(y, lags, options);
```

By default, optimization is performed unconstrained and `cminwel.m` is used (this is `options.max_compute = 3`). The following options can be set in the maximization step:

- `options.lb` and `options.ub` set the lower and upper bounds for the optimization. Both are row array vectors of the same size of `options.priors.tau`.
- `options.max_compute` is a scalar, selecting the maximization routine to be employed.

The options are:

- `options.max_compute = 1` uses the MATLAB `fminunc.m` (unconstrained)
- `options.max_compute = 2` uses the MATLAB `fmincon.m` (constrained)
- `options.max_compute = 3` uses the Chris Sims's `cminwel.m`
- `options.max_compute = 7` uses the MATLAB `fminsearch.m`

1, 2, 3 doesn't typically work, so start with 7 (it does not use derivatives).

The first three are derivative-based algorithms; the latter is a direct search (simplex) method based on function comparisons. While typically slower, it is useful in situations where derivatives are not well behaved.

# Bayesian Local projections III

The output of the DM function is a field with a number of sub-fields. The most relevant ones are

- `dm.ir_blp` is a  $(n \times \text{hor} \times n \times K)$  matrix containing Bayesian local projections impulse responses obtained obtained with recursive identification. The first dimension corresponds to the variables, the second to the horizon, the third to the disturbances, and the fourth to the responses obtained with particular draw from the posterior distribution of the local projection parameters.
- `dm.irproxy_blp` is a  $(n \times \text{hor} \times 1 \times K)$  matrix containing Bayesian local projection impulse responses obtained with IV identification. The first dimension corresponds to the variables, the second to the horizon, the third to the shock, and the fourth to the responses obtained with particular draw from the posterior distribution of the local projection parameters.
- `dm.bforecasts.no_shocks` is a  $(\text{hor} \times n \times K)$  matrix containing unconditional forecasts. The first dimension corresponds the horizon, the second to the variable, and the third to the draw from the posterior distribution of the parameters. These trajectories are constructed assuming that all future shocks are zeros (forecasts only account for parameter uncertainty).
- `dm.bforecasts.with_shocks` is a  $(\text{hor} \times n \times K)$  matrix containing unconditional forecasts with shocks. The first dimension corresponds the horizon, the second to the variable, and the third to the draw from the posterior distribution of the parameters. These trajectories account for uncertainty in the parameter estimates and in the shocks realization.

# Practice VII

Sample codes are in ... \BVAR\examples\BVAR tutorial.

- Bayesian and Classical local projection (example\_7\_LP.m) .
- Forecasting using local projections (example\_7\_LP.m).

# Outline

- 1 Introduction
- 2 Turning point dating
- 3 Trend and cycle decompositions
- 4 Classical VAR Estimation and Inference
- 5 Bayesian VARs Estimation and Inference
- 6 Impulse response functions (IRF)
- 7 Identification restrictions
- 8 Forecast error and historical decompositions
- 9 Time variations
- 10 Large scale VARs, FAVARs, and DFM
- 11 Panels of VARs
- 12 Mixed Frequency VARs
- 13 Forecasts
- 14 Local projections and Direct forecasts
- 15 Appendix**

# Some notation I

- **VAR( $p$ )**: A vector of autoregression with  $p$  lags:

$$y_t = \Phi_1 y_{t-1} + \dots + \Phi_p y_{t-p} + \Phi_0 + u_t \quad u_t \sim N(0, \Sigma) \quad (1)$$

where  $y_t$  is  $n \times 1$  vector,  $\Phi_j$  suitable matrices

- **SURE**: A VAR( $p$ ) be rewritten as a SURE:

$$Y = X\Phi + E$$

- **Forecasts and IRF**: Any VAR( $p$ ) can be expressed as a (companion form) VAR(1)

$$x_t = Fx_{t-1} + F_0 + Gu_t$$

- **Decompositions**: Under some restrictions a VAR( $p$ ) can be written as a Vector MA (VMA( $\infty$ ))

$$y_t = u_t + \Psi_1 u_{t-1} + \dots + \Psi_t u_1 + \bar{\Psi}_t$$

where  $\Psi_j$  for  $j = 1, \dots, t$  and  $\bar{\Psi}_t$  are functions of  $(\Phi_1, \dots, \Phi_p)$ .

# Some notation II

- $LP(h, d, q)$ : Local projection at horizon  $h$  is given by

$$y_{it+h} = \Phi_1 y_{it-1} + \dots + \Phi_p y_{it-d} + \Lambda_1 y_{jt-1} + \dots + \Lambda_q y_{jt-q} + \Phi_0 + \zeta v_t + u_{it+h} \quad (2)$$

where  $y_{it}$  is scalar  $i \neq j$ ,  $\Phi_j, \Lambda_j$  are suitable matrices,  $v_t$  is a shock of interest,  $u_{it+h} \sim iid(0, \Sigma)$ . Under certain conditions:  $\Psi_j = \zeta \Phi_j$ .

- **Trend and Cycle decompositions**: Any observable scalar (vector)  $y_t$  be decomposed into two latent components

$$y_t = \tau_t + c_t$$

with suitable restrictions on  $\tau_t, c_t$  or both.



# Empirical and structural innovations

- Reduced-form VAR innovations  $u_t$  and structural disturbances  $\nu_t$  relationship:

$$u_t = \Omega \nu_t \equiv \Omega_0 Q \nu_t$$

where  $E(\nu_t \nu_t') = I$  and  $\Omega \Omega' = \Sigma$ ,

- $\Omega_0$  is (typically) the Cholesky decomposition of  $\Sigma$  and  $Q$  is an orthonormal rotation such that  $Q'Q = QQ' = I_n$ .
- To recover  $\nu_t$ , we need to impose restrictions on  $\Omega$ .
- This is because  $\Sigma$  only contains  $n(n+1)/2$  estimated elements, while  $\Omega$  has  $n^2$  elements.